

Motivation

This work aims to bridge the gap between the prerequisite large-scale training and the constraint of data sharing due to confidentiality. We propose the Unidirectional Split Learning with Contrastive Loss (UniCon) to tackle Visual Question Answering (VQA) tasks for distributed data silos. UniCon trains a global model by transferring knowledge from different clients' local tasks learning refined cross-modal representations via contrastive learning. This work is the first study of VQA under the constraint of data confidentiality using self-supervised learning.

Introduction

- The real-world deployment of VQA in safety-critical applications such as healthcare needs to address the robust architecture design. Previous studies do not touch on the privacy of VQA with multi-modality data.
- The semantic notions of answers are usually not well correlated with the inputs reducing the generality of the trained model to unknown samples.

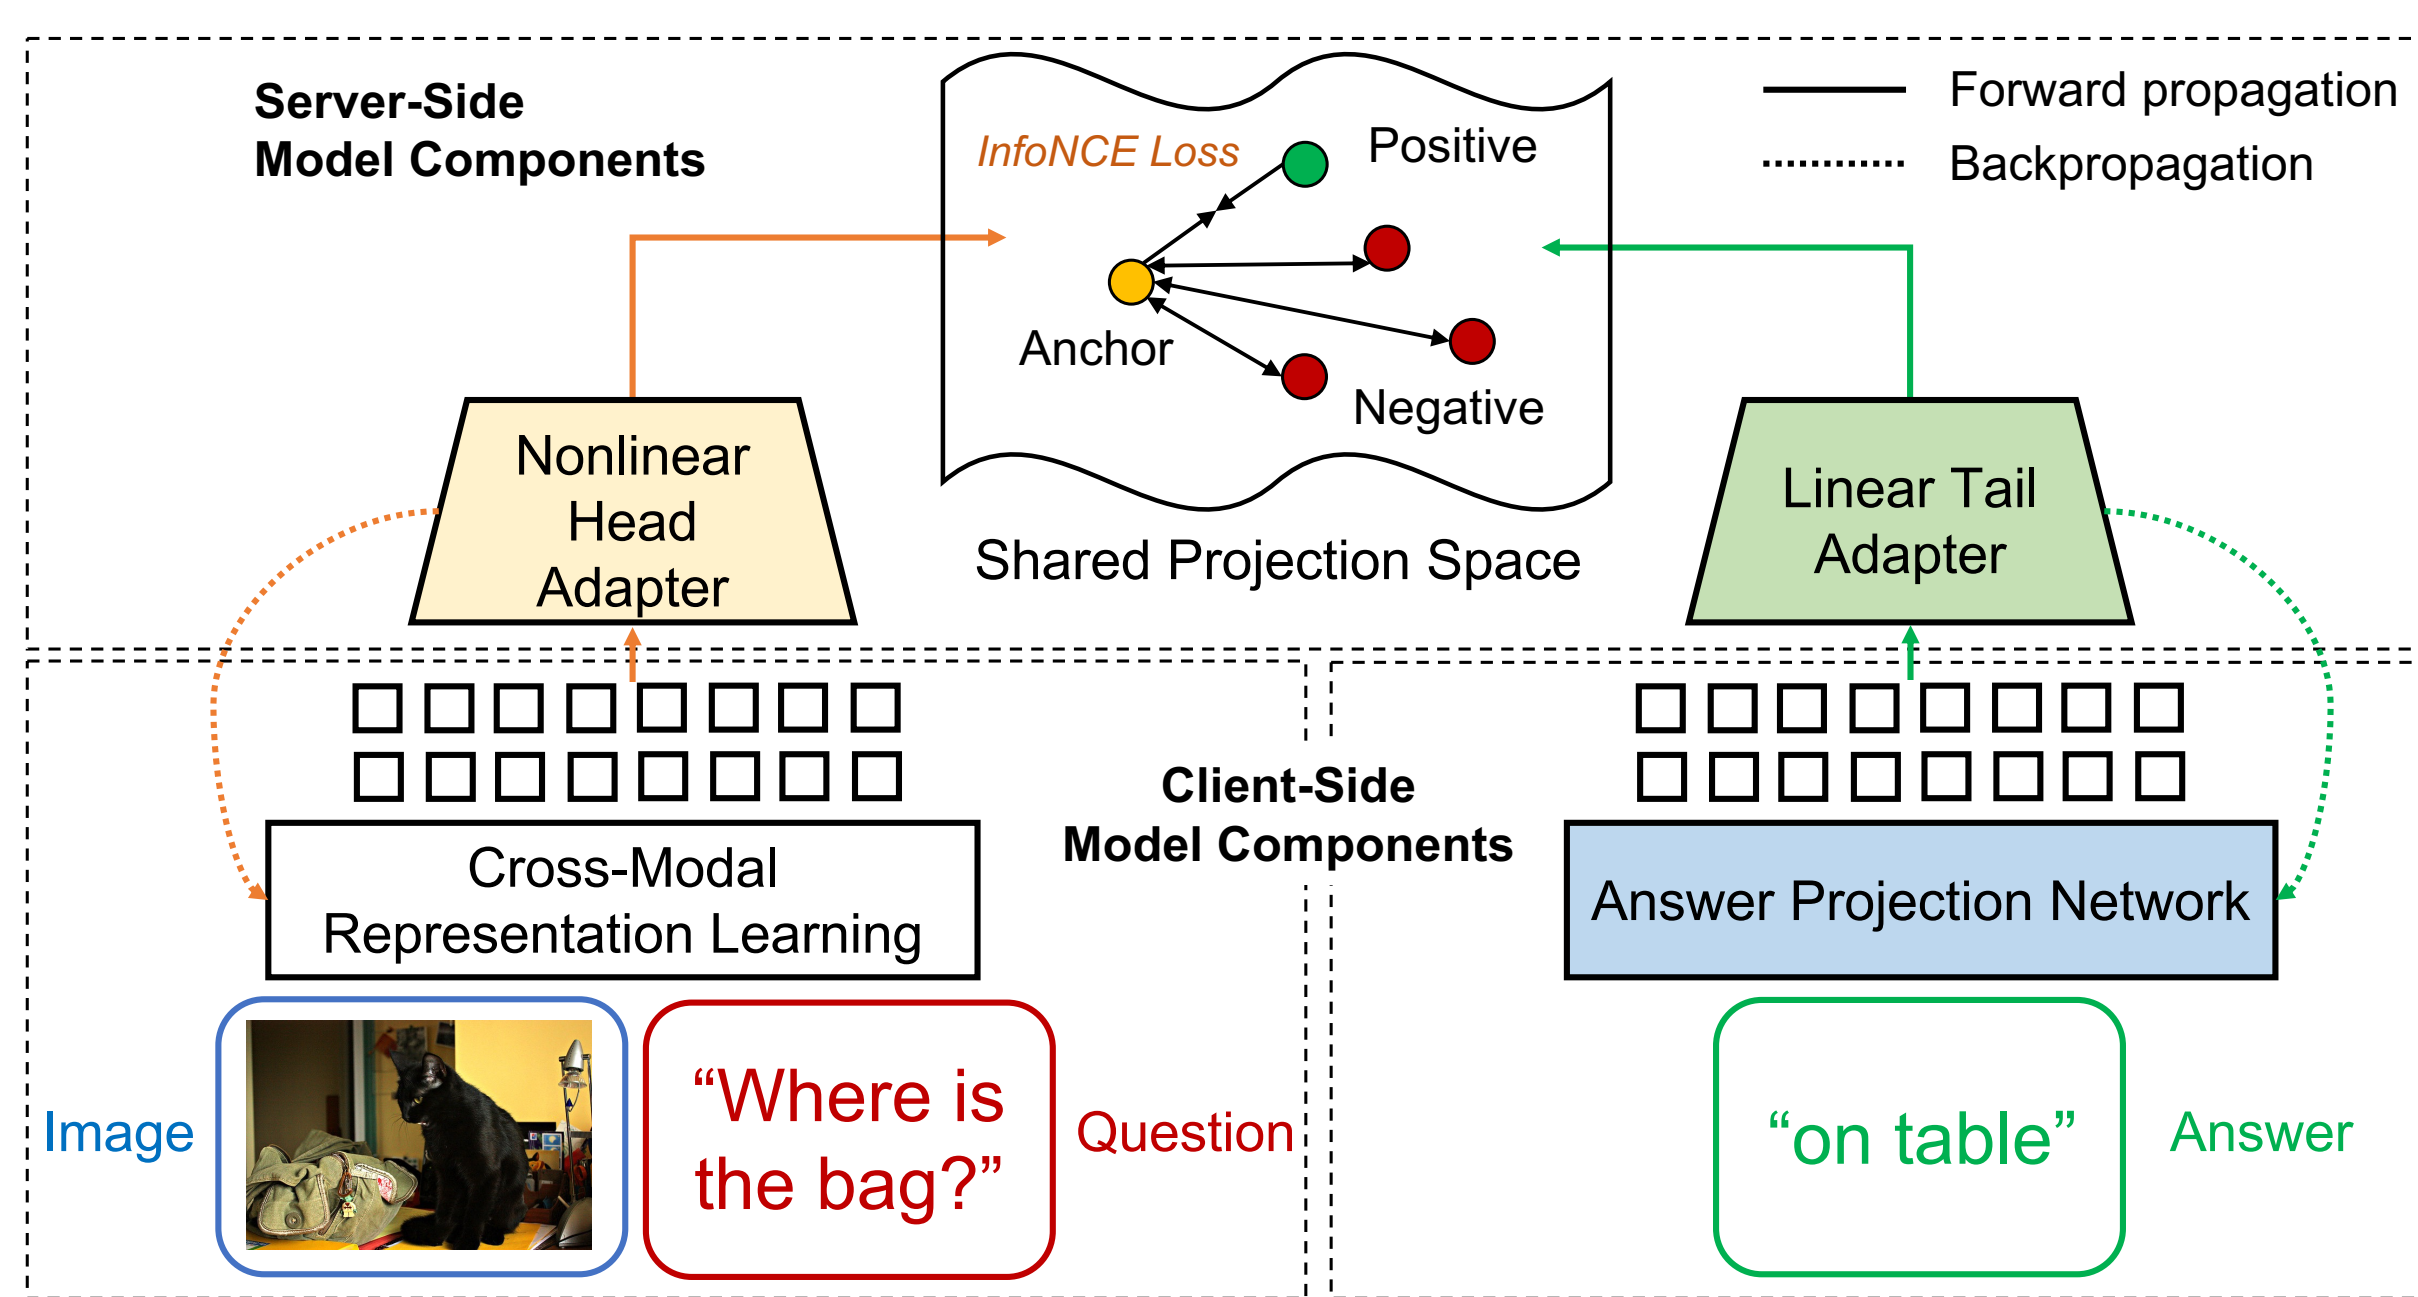


Figure 1: The overall architecture of Unidirectional Split Learning with Contrastive Loss (UniCon).

Methodology

Split Learning for Visual Question Answering

- The goal of the VQA model is to predict the correct answer y given the input pair of an image and a question (x_i, x_q) .
- Split Learning (SL) guarantees the confidentiality of a client by splitting a complete model into different components, such that the input data and the complete model are not shared during the training.




Figure 2: The comparison between Split Learning and UniCon.

Unidirectional Split Learning with Contrastive Loss

We propose the use of the contrastive loss in Split Learning to correlate vision contents and language semantic notions such that each model component learns better-refined cross-modal representations for the VQA tasks. Moreover, the direction of the propagation is unidirectional in UniCon, which allows the components to be computed simultaneously.

- The Answer Projection Network (APN) f_{APN} embeds the answer language contexts y into a feature vector $v_{APN} \in \mathbb{R}^P$.

- We propose the use of two adapter networks to project the outputs from different model components into the shared projection space.
- We replace a VQA model's output layer with the Nonlinear Head Adapter (NHA) network f_{NHA} that projects the high-level cross-modal representations $v_{NHA} \in \mathbb{R}^S$. The Linear Tail Adapter (LTA) projects the low-level representations v_{APN} of APN into the shared projection space $v_{LTA} \in \mathbb{R}^S$.

Learning with Information Noise Contrastive Estimation loss

- The relevant NHA and LTA outputs of the same input triplets within one training batch are employed as positive pairs. $\{(v_{NHA,i}, v_{LTA,i})\}_{i=1}^B$ where B is the batch size.
- Given a NHA output $v_{NHA,i}$, any irrelevant LTA outputs $\{v_{LTA,j} | j \neq i\}_{j=1}^B$ are employed as the negative keys of the NHA output.
- Train the model by aligning the knowledge between the component outputs in positive pairs while discouraging the similarity between the outputs in negative pairs.

$$\mathcal{L} = - \sum_{i=1}^B \log \frac{\exp(v_{NHA,i} \cdot v_{LTA,i} / \tau)}{\sum_{j=1}^B 1_{[j \neq i]} \exp(v_{NHA,i} \cdot v_{LTA,j} / \tau)}.$$

Results

- Extensive experiments with seven state-of-the-art VQA models.
- Contrastive learning-based models are trained on the entire VQA dataset.
- UniCon was trained on the two complementary subsets of the entire dataset facilitating model training by reusing knowledge of different local tasks, without disclosing client data and models.

VQA Models	Contrastive learning-based VQA				UniCon (%)			
	Overall	Yes/No	Number	Other	Overall	Yes/No	Number	Other
BAN	36.23	66.90	12.71	19.11	35.11	63.84	11.06	19.61
BUTD	45.08	75.82	29.27	25.86	40.96	66.98	13.34	28.74
MFB	46.98	73.95	32.81	30.20	42.43	68.65	23.33	27.52
MCAN-s	53.18	81.06	41.95	34.93	48.42	74.93	30.88	32.89
MCAN-l	53.32	81.21	42.66	34.90	48.44	77.44	30.72	32.01
MMNas-s	51.54	78.06	39.76	34.46	45.14	70.55	28.04	30.33
MMNas-l	53.82	80.06	42.86	36.75	49.89	74.85	36.88	34.33

Table 1: Performance evaluation for different models.

We considered different client numbers of $K = \{1, 2, 3, 4, 5\}$. The k th client has a sample size of $N^{(k)}$. $\sum_{k=1}^K N^{(k)} = N$, where N is the entire training data size.

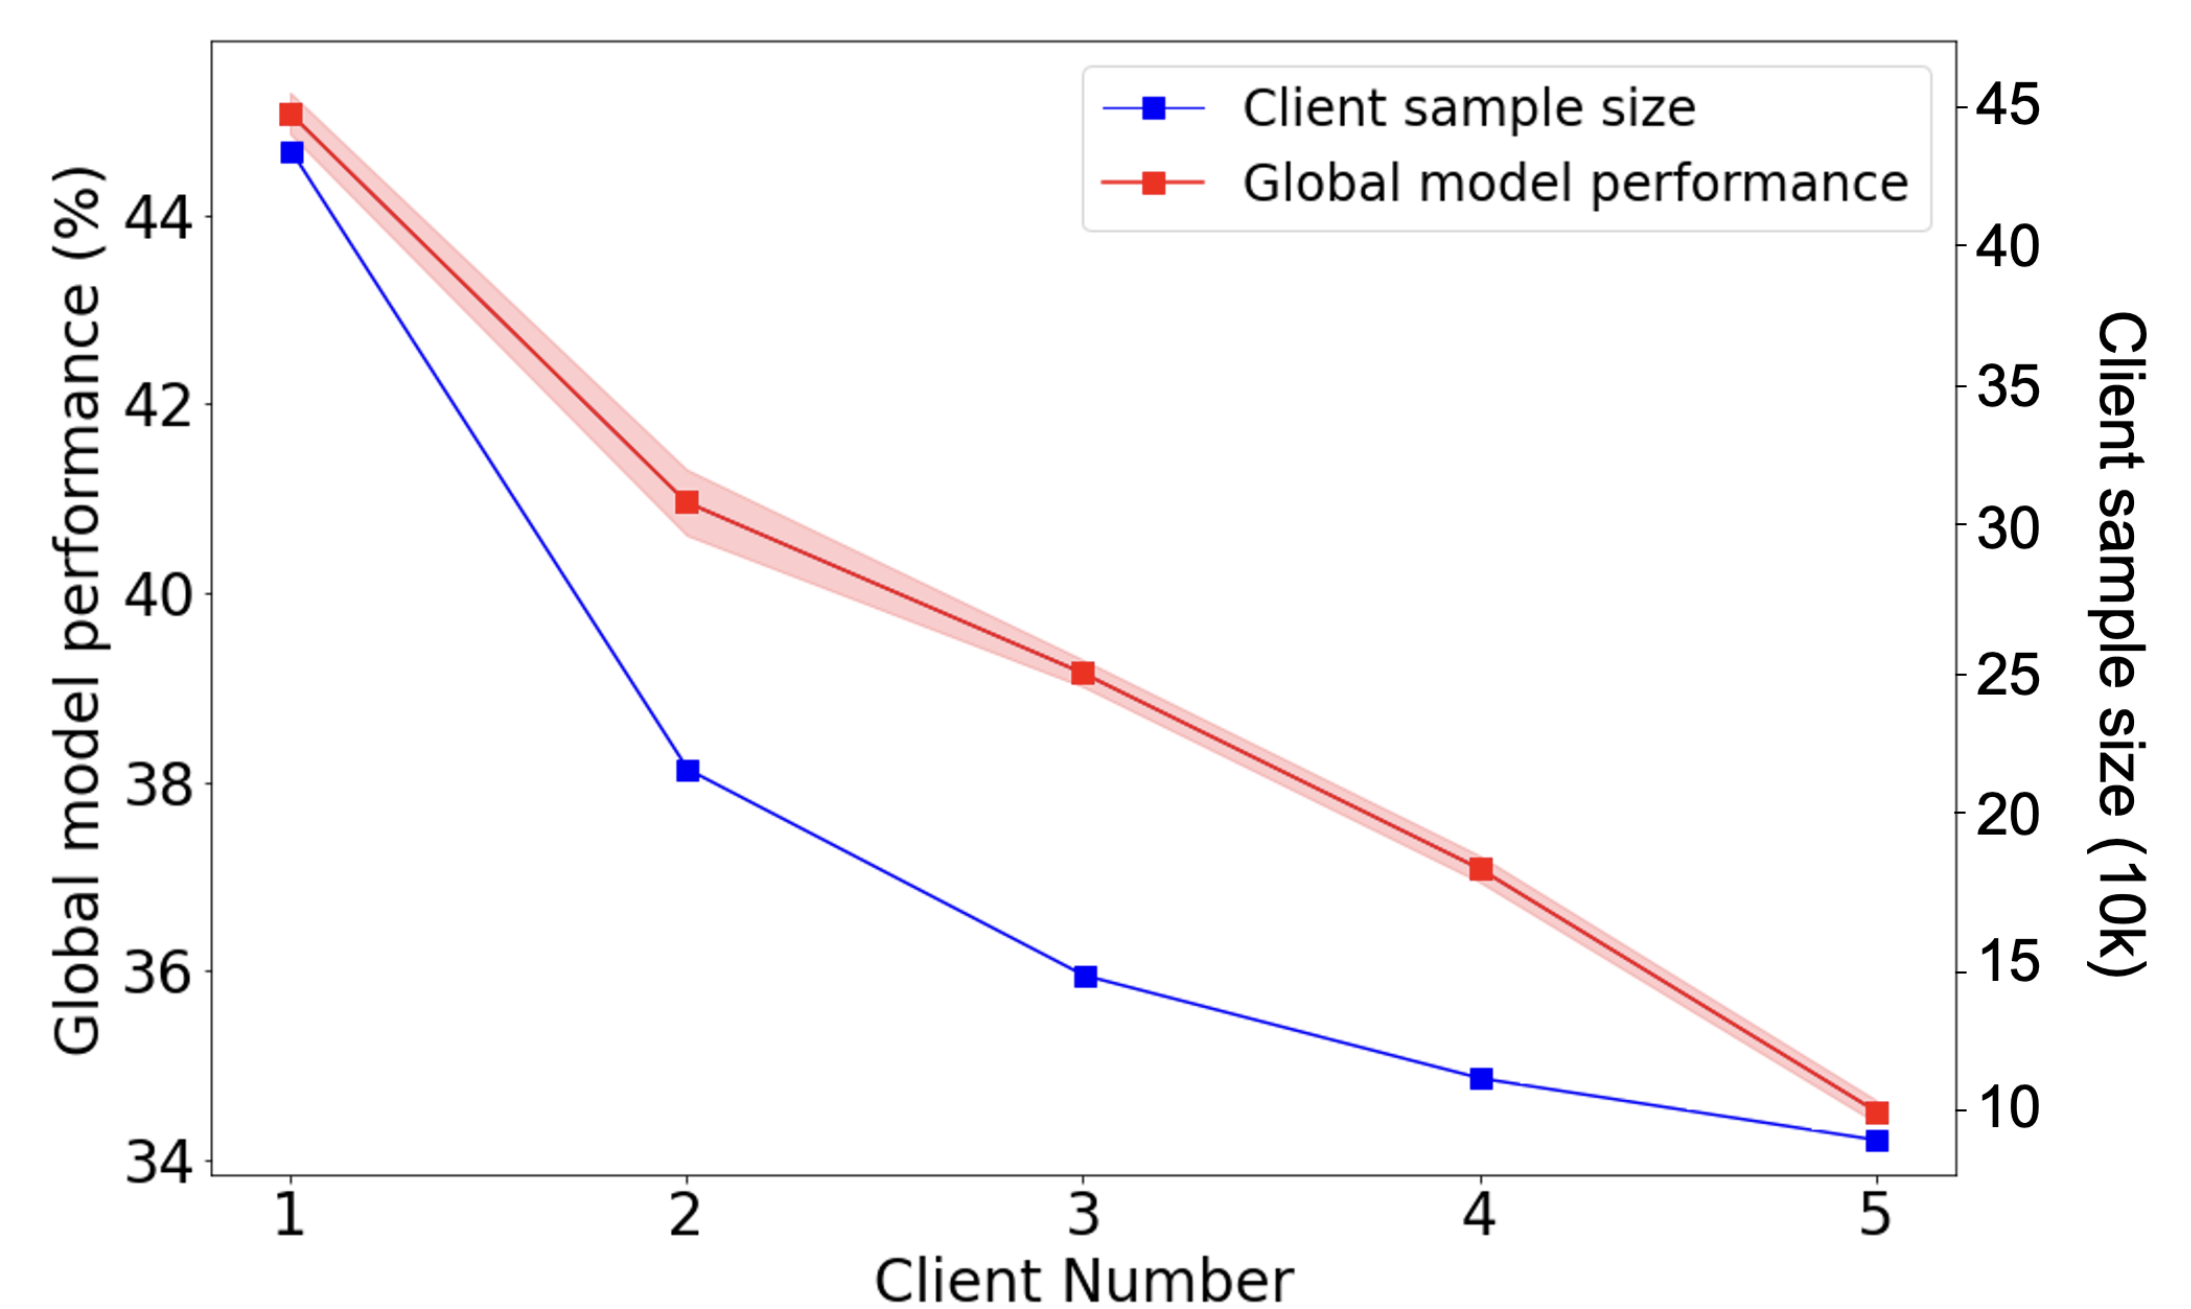


Figure 3: The tradeoff between the client number and the model performance.

Conclusion

We proposed the Unidirectional Split Learning with Contrastive loss (UniCon) to tackle the robust learning of VQA under the constraint of data confidentiality. UniCon learns refined cross-modal representations by aligning model component outputs based on contrastive learning. We presented an in-depth evaluation on the VQA-v2 dataset using a wide range of state-of-the-art models. The results showed that UniCon can be applied in various learning settings.